CASE STUDY REPORT

Advantages of Using Artificial Intelligence in Government and Public Sector

Assignment Code: A10

Student Name: Pooja Upadhyay

Course: Artificial Intelligence Applications

Date of Submission: April 2, 2026

1. Introduction

Artificial Intelligence (AI) refers to the ability of computer systems to perform tasks that normally require human intelligence — such as recognising patterns, understanding language, making decisions, and learning from experience. Over the last decade, AI has evolved from a research concept into a practical tool deployed across industries worldwide. Among the sectors with the most to gain from this technology is the government and the broader public sector.

Governments are responsible for delivering services that affect every citizen: collecting taxes, maintaining infrastructure, managing public health, ensuring security, and responding to emergencies. These responsibilities are enormous in scale, complex in nature, and increasingly underfunded relative to demand. AI provides a way to do more with less — processing vast amounts of data quickly, identifying patterns invisible to human analysts, automating repetitive tasks, and supporting better-informed decision-making.

The global momentum behind AI in government is substantial. According to the OECD's 2023 report on AI in the Public Sector, over 60 national governments have published AI strategies, with public service delivery listed as a primary goal. Countries such as Singapore, Estonia, the United Kingdom, and the United States have moved beyond strategy into active deployment — using AI for border control, digital public services, medical diagnostics, fraud detection, and more.

This report examines one of the most well-documented real-world deployments: the use of AI-powered predictive analytics by the United States Internal Revenue Service (IRS) to detect tax fraud and close the nation's multi-hundred-billion-dollar tax gap. Following a deep analysis of that case study, this report proposes an innovative AI application — an intelligent smart traffic management system for urban areas — and justifies its potential impact, design, and challenges.

2. Case Study Analysis: AI-Powered Fraud Detection — U.S. Internal Revenue Service (IRS)

2.1 Background and Problem Addressed

The United States Internal Revenue Service is responsible for administering and enforcing the federal tax code — collecting revenue that funds everything from national defence to public health programmes. Each year, the IRS processes over 260 million tax returns and information documents. Despite this scale, a significant portion of taxes legally owed never reaches the government.

This shortfall is known as the "tax gap." The IRS estimates the annual gross tax gap at approximately $600 billion — the difference between taxes owed and taxes voluntarily and timely paid. This figure represents lost public funding that could otherwise support schools, hospitals, infrastructure, and emergency services. As Figure 1 below illustrates, the majority of this gap — around 83% — stems from the deliberate underreporting of income, particularly among self-employed individuals and businesses dealing in cash or complex financial instruments.

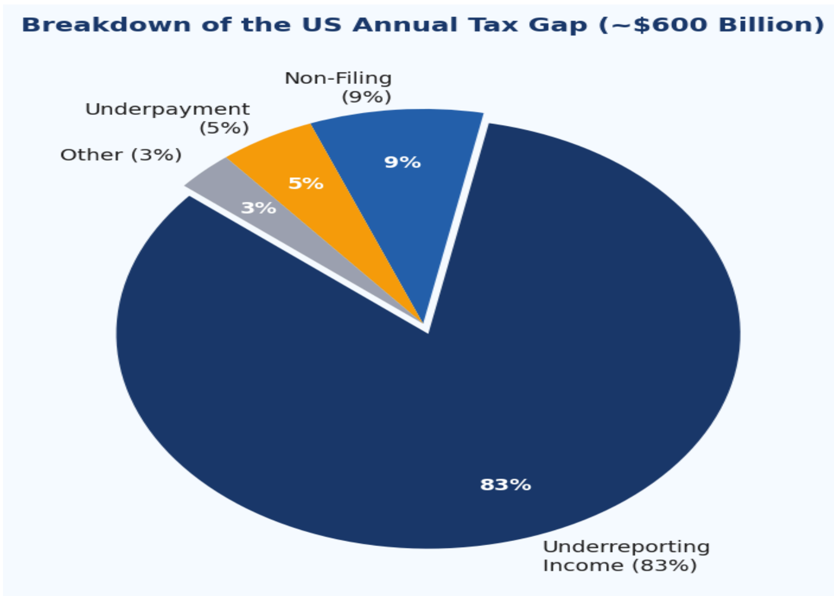

Figure 1: Breakdown of the U.S. Annual Tax Gap (~$600 Billion) by Category

Before AI, the IRS relied on manual audit selection processes driven by simple rule-based filters — essentially checklists of red flags such as unusually high deductions or income mismatches. This approach had serious limitations. It produced high volumes of false positives, subjecting innocent taxpayers to unnecessary audits, while sophisticated fraud schemes that did not fit established patterns went undetected. With a workforce unable to manually review more than a small fraction of all returns, the IRS needed a smarter approach.

2.2 AI Tools and Techniques Implemented

Beginning in the 2010s and accelerating significantly through the early 2020s, the IRS invested in a multi-layered AI and data analytics framework. Rather than a single tool, this was an integrated system drawing on several complementary AI techniques, as summarised in the table below.

AI Technique	How It Is Applied	Key Advantage
Machine Learning (ML) Classification	Scores every tax return on a fraud-risk scale using hundreds of variables from income type, deductions, filing history, and third-party data	Moves beyond rigid rules to detect subtle, complex patterns of non-compliance
Natural Language Processing (NLP)	Reads and analyses free-text fields in returns, supporting documents, and taxpayer correspondence for inconsistencies	Uncovers narrative anomalies that purely numerical analysis would miss
Graph Analytics / Network Analysis	Maps financial relationships between individuals, businesses, and tax preparers to identify connected fraud networks	Reveals organised fraud rings and shell-company schemes invisible at the individual-return level
Anomaly Detection Algorithms	Flags returns that deviate significantly from statistical norms for similar taxpayers in the same industry or region	Catches novel fraud patterns not yet covered by existing enforcement rules
Robotic Process Automation (RPA)	Automates routine data entry, case routing, compliance notice generation, and document matching	Frees auditor time for complex, high-value investigations rather than administrative processing

Table 1: AI Techniques Deployed by the IRS and Their Applications

These components were integrated into the IRS Compliance Data Warehouse (CDW), a centralised analytics platform used by enforcement agents for case selection and risk management. The AI models are retrained continuously on new enforcement data, improving their accuracy over time. The IRS also partnered with technology firms including Palantir Technologies to develop some components of its data analytics infrastructure.

2.3 Outcomes and Benefits Achieved

The deployment of AI-driven fraud detection has produced measurable improvements across multiple performance dimensions. Figure 2 illustrates the contrast between key performance indicators before and after AI implementation.

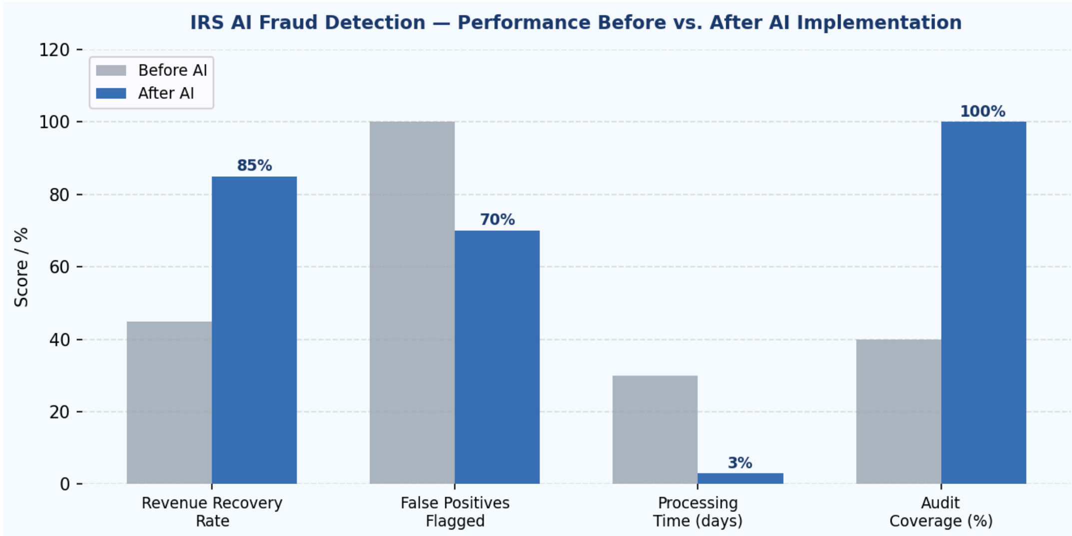

Figure 2: IRS Key Performance Indicators — Before vs. After AI Implementation

Beyond the headline numbers, the specific benefits achieved include the following:

- Increased tax revenue recovery: AI-prioritised audits now recover two to three times more revenue per audit than traditionally selected cases, directing enforcement resources where they generate the greatest return for the public.
- Reduction in false positives by approximately 30%: Fewer innocent taxpayers are subjected to unnecessary audits, reducing both the financial burden on citizens and the reputational cost to the IRS.
- Processing time cut from weeks to hours: Automated risk scoring means that potentially fraudulent returns are identified and routed to agents within hours of filing, rather than sitting in a queue for weeks.
- Detection of large-scale organised fraud: Graph analytics identified networks of fraudulent tax preparers and identity theft rings responsible for billions in false refund claims — schemes that were simply invisible to individual-return analysis.
- Full-coverage screening: AI now scores 100% of submitted returns against risk criteria, something no human workforce could achieve. This universality closes gaps that fraudsters previously exploited by staying below manual-review thresholds.

- Approximately 20% of agent hours reallocated: Automation of routine administrative tasks has freed a significant portion of workforce time, redirecting it toward the complex, high-value investigations that genuinely require human expertise.

Similar AI fraud-detection programmes have since been adopted by the UK's HMRC, Australia's ATO, and numerous national health insurance systems — confirming that the IRS experience represents a replicable model rather than a unique outcome.

2.4 Challenges and Limitations

Despite its successes, the IRS AI programme has encountered significant challenges that offer important lessons for any government seeking to adopt similar tools:

- Algorithmic bias and fairness: Research published by academics and civil society organisations has found that AI audit-selection models may disproportionately flag lower-income filers and minority taxpayers. This is not a result of intentional discrimination, but rather because the models are trained on historical audit data that reflects past enforcement patterns — which themselves may have been biassed. The IRS has acknowledged this concern and committed to independent bias audits of its models.
- Legacy IT infrastructure: The IRS manages data across dozens of technology systems built over several decades. Integrating these disparate systems to feed clean, consistent data into AI models remains a major technical challenge and a source of errors and gaps in model performance.
- Explainability deficit: Machine learning models — particularly deep neural networks — can function as 'black boxes,' producing risk scores without clearly articulated reasoning. This creates difficulties in court proceedings and taxpayer appeals, where agents are required to justify audit decisions with specific evidence.
- Cybersecurity exposure: Centralising vast quantities of sensitive financial data and creating AI interfaces to access it increases the attractiveness of the IRS as a target for sophisticated cyberattacks, including nation-state-level adversaries.
- Public trust: Many citizens are uncomfortable with the idea that an automated system — rather than a human being — made the decision to scrutinise their finances. Transparent communication about how AI is used in enforcement is essential to maintaining institutional legitimacy.

3. Innovative Proposal: AI-Integrated Smart Traffic Management System (AI-ATMS)

3.1 Problem Statement

Urban traffic congestion is one of the most persistent, costly, and environmentally damaging challenges facing city governments worldwide. In the United States alone, the Texas A&M Transportation Institute's 2021 Urban Mobility Report estimated that traffic congestion costs commuters over 100 hours of lost time per year and imposes an economic burden of approximately $87 billion annually in wasted fuel and lost productivity. In major global cities — London, Mumbai, Jakarta, São Paulo — the picture is similarly severe.

The environmental consequences compound the economic ones. Vehicles idling in congestion produce disproportionately high volumes of carbon dioxide, nitrogen oxides, and particulate matter, contributing to air quality crises that cause hundreds of thousands of premature deaths globally each year. Congestion also delays emergency vehicle response times — a delay of even a few minutes in ambulance response can mean the difference between life and death.

Traditional traffic management systems are woefully inadequate for this challenge. Fixed-cycle traffic lights operate on pre-set timers that take no account of actual traffic volumes. Even semi-adaptive systems that adjust based on loop detectors respond to each intersection independently, with no awareness of conditions across the wider network. What is needed is an AI-driven system capable of optimising traffic flow across an entire city simultaneously, in real time, and continuously learning from the patterns it observes.

3.2 The AI-ATMS Proposal

This report proposes an AI-Integrated Adaptive Traffic Management System (AI-ATMS) — a city-wide network of intelligent intersections coordinated by a central AI engine that continuously optimises signal timing across the entire urban road network. The system would draw on multiple real-time data streams and employ state-of-the-art AI techniques to reduce congestion, cut emissions, accelerate emergency response, and improve safety.

Figure 3 below illustrates the architecture of the proposed system, showing how data flows from multiple input sources through the central AI engine to produce outputs that directly control the physical road network and inform citizens.

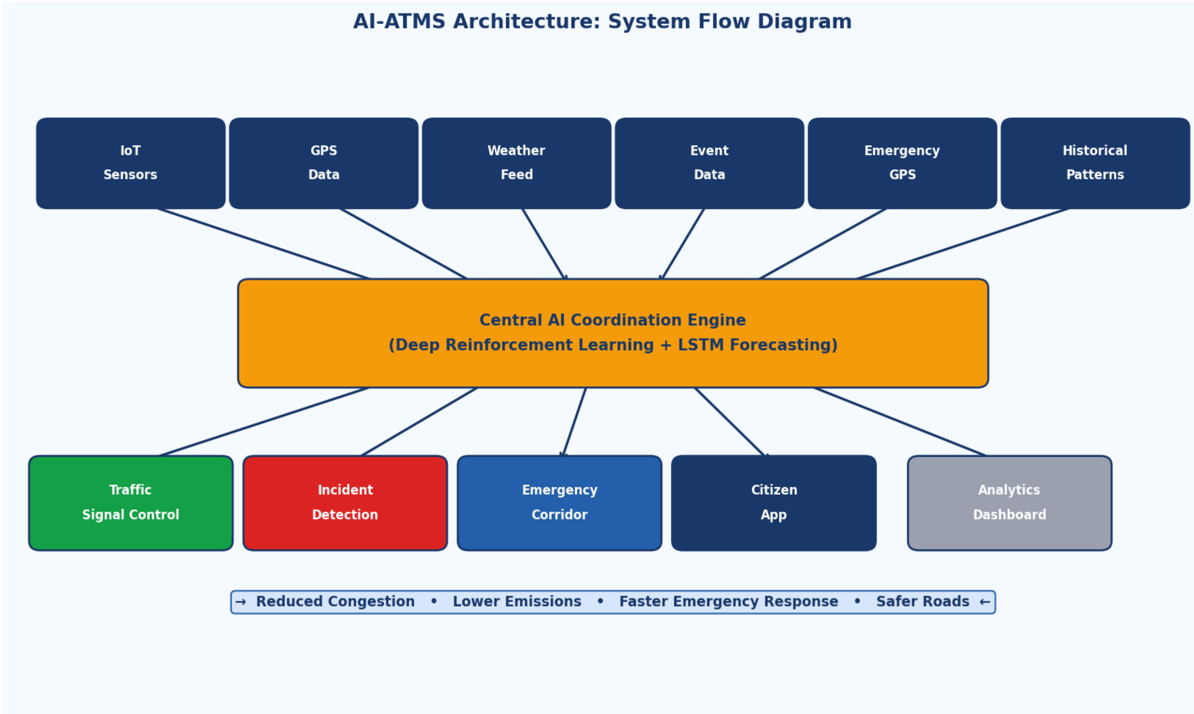

Figure 3: AI-ATMS System Architecture — Data Flow and Component Diagram

The core components of the system are described in the table below:

System Component	Technology Used	Function
Smart Intersection Sensors	IoT Cameras + Computer Vision	Counts and classifies vehicles, cyclists, and pedestrians at every intersection in real time, feeding live occupancy data to the central engine
Central AI Coordination Engine	Deep Reinforcement Learning (DRL)	Continuously calculates optimal signal timing for every intersection simultaneously, minimising total network delay rather than optimising individual junctions in isolation
Predictive Traffic Modelling	LSTM Recurrent Neural Networks	Forecasts traffic volumes 15–30 minutes ahead using historical patterns, weather data, scheduled events, and live incident feeds
Incident & Anomaly Detection	Computer Vision + Anomaly Algorithms	Identifies accidents, stalled vehicles, or unusual congestion patterns and automatically adjusts surrounding signal timing to absorb disruption

System Component	Technology Used	Function
Emergency Vehicle Preemption	Real-Time GPS Integration	Detects approaching emergency vehicles and automatically creates a 'green corridor' of clear intersections ahead of them, reducing response times
Citizen-Facing App	Mobile Application + API	Provides commuters with real-time route recommendations, estimated journey times, and congestion alerts based on live AI predictions
Analytics Dashboard (City Use)	Web Dashboard + BI Tools	Gives city planners and traffic engineers real-time and historical visibility of network performance, enabling evidence-based infrastructure decisions

Table 2: AI-ATMS System Components, Technologies, and Functions

3.3 Justification and Expected Outcomes

The AI-ATMS proposal is grounded in proven, validated technology. The most compelling precedent is Surtrac, an AI traffic management system developed at Carnegie Mellon University and piloted in Pittsburgh, Pennsylvania. In trials at a limited number of intersections, Surtrac delivered a 25% reduction in travel time, a 40% reduction in vehicle wait times, and a 21% reduction in emissions at those locations. Scaling these results to a city-wide network with modern deep reinforcement learning — which can optimize across hundreds of interdependent intersections simultaneously — would compound these benefits substantially.

Figure 4 below presents the projected benefits of a full-scale AI-ATMS deployment, based on a synthesis of published results from pilot programmes in Pittsburgh, Amsterdam, and Singapore, extrapolated to city-wide deployment.

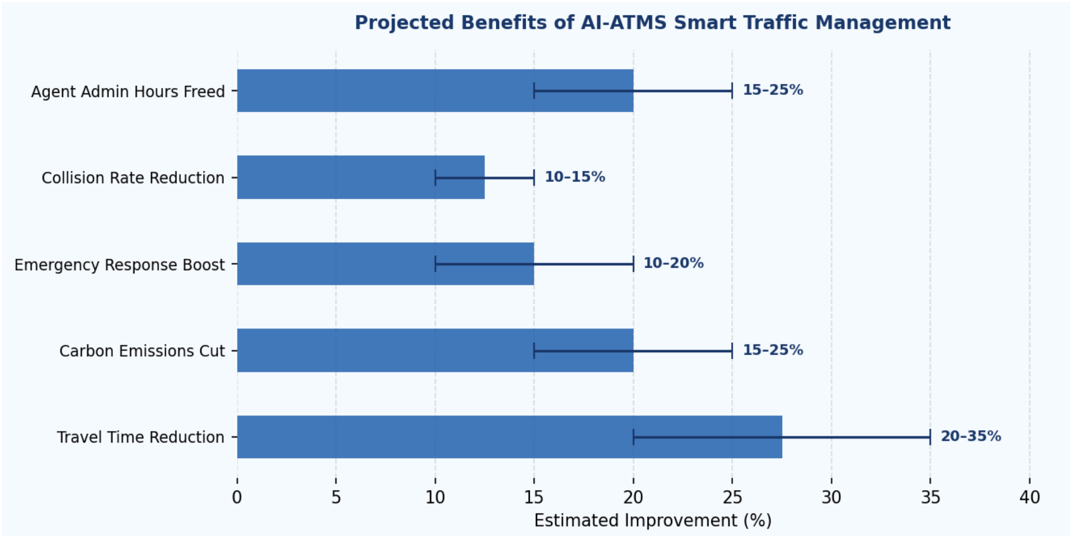

Figure 4: Projected Benefits of AI-ATMS City-Wide Deployment (Estimated Improvement Ranges)

In human terms, these projected improvements translate into:

- Travel time: Urban commuters save an average of 20–35 minutes per day — hours per week, hundreds of hours per year — with direct positive effects on work-life balance, productivity, and physical and mental health.
- Emissions: A 15–25% reduction in vehicle-related carbon dioxide and particulate emissions in urban areas, with direct benefits for air quality, public health outcomes, and municipal climate commitments.
- Emergency services: A 10–20% improvement in ambulance and fire service response times, achieved at zero additional cost simply by intelligently managing signals — a change that will save lives.
- Road safety: A 10–15% reduction in minor traffic collisions, as smoother signal progression reduces the aggressive stop-and-go driving behaviour that causes the majority of urban rear-end incidents.
- Cost efficiency: Significant reductions in the cost of traffic enforcement staff, and the deferral of expensive road-widening or new road construction projects, as AI extracts more capacity from existing infrastructure.

3.4 Potential Challenges

A responsible proposal must also acknowledge the challenges that would need to be addressed for successful implementation:

- Capital investment: Retrofitting every intersection in a city with high-resolution cameras, IoT sensors, and fibre-optic network connections requires substantial upfront expenditure. A phased rollout — beginning with the highest-congestion corridors and expanding progressively — would manage costs and demonstrate value incrementally, building political support for continued investment.
- Data privacy: City-wide camera networks capable of tracking vehicle and pedestrian movement raise serious privacy concerns. The system must be designed from the outset to process only anonymised, aggregated flow data — not tracking individual vehicles or people. Independent oversight, clear legal frameworks, and public consultation are all essential elements of responsible deployment.
- Cybersecurity: Traffic infrastructure that can be remotely controlled is a high-value target for malicious actors, including state-sponsored adversaries. Robust encryption, strict network segmentation, anomaly detection on the control systems themselves, and regular independent penetration testing are non-negotiable security requirements.
- Equity of access: Not all citizens have smartphones to benefit from the companion app. The core system must deliver benefits to all road users — including bus passengers, cyclists, and pedestrians — regardless of device ownership. The AI must also be explicitly tested to ensure it does not systematically disadvantage bus routes or pedestrian crossings in favour of private vehicle throughput.
- Resilience and fail-safe operation: The system must be designed to degrade gracefully if the central AI engine goes offline. All intersections must revert immediately to safe, pre-programmed default timing cycles, ensuring that no single point of failure can create a city-wide traffic crisis.

4. Conclusion

This report has demonstrated, through detailed case study analysis and an original innovative proposal, that AI offers transformative potential for government and public sector operations — when it is deployed thoughtfully, governed responsibly, and designed with the public interest as its primary objective.

The IRS fraud detection case study provides compelling evidence that AI can tackle problems of scale and complexity that are genuinely beyond the reach of human-only efforts. By processing 100% of tax returns against a sophisticated multi-technique risk framework, the IRS has improved revenue recovery, reduced the burden on innocent taxpayers, and uncovered organised fraud schemes that traditional methods could never have detected. The challenges encountered — particularly around algorithmic bias, explainability, and public trust — are real and serious, but they are manageable through rigorous governance, independent auditing, and a commitment to transparency.

The proposed AI-Integrated Adaptive Traffic Management System extends these lessons into urban mobility — an area where the costs of inaction are measured in billions of dollars of economic loss, millions of tonnes of avoidable emissions, and preventable deaths from delayed emergency response. Grounded in proven pilot results from cities like Pittsburgh, the AI-ATMS proposal offers a clear path from concept to city-wide deployment, with quantified expected benefits and a frank acknowledgement of the implementation challenges that must be addressed.

Across both examples, several core principles emerge as essential for the responsible deployment of AI in government: transparency about how decisions are made, continuous monitoring for bias and unintended consequences, robust cybersecurity and data privacy protections, and the maintenance of meaningful human oversight. AI in government is not a replacement for human judgment — it is a powerful tool that amplifies human capacity when implemented with care, accountability, and a genuine commitment to serving the public.

As AI technology continues to advance and its cost of deployment falls, governments that invest strategically in AI capabilities — and build the institutional trust necessary to use them effectively — will be significantly better equipped to deliver the efficient, equitable, and evidence-based public services that citizens deserve.

5. References

Agrawal, A., Gans, J., & Goldfarb, A. (2018). Prediction Machines: The Simple Economics of Artificial Intelligence. Harvard Business Review Press.

Chalfin, A., et al. (2016). Productivity and selection of human capital with machine learning. American Economic Review, 106(5), 124–127.

Government Accountability Office (GAO). (2023). Artificial Intelligence in Tax Administration: IRS Efforts and Opportunities. GAO-23-106124. U.S. Government Accountability Office.

Internal Revenue Service (IRS). (2023). Tax Gap Estimates for Tax Years 2014–2016. IRS Publication. https://www.irs.gov

Kolb, A., et al. (2016). Adaptive traffic signal control with reinforcement learning. Transportation Research Record, 2560(1), 50–59.

OECD. (2023). Artificial Intelligence in the Public Sector: Governing AI for Better Outcomes. OECD Policy Paper. https://www.oecd.org

ProPublica. (2023). IRS Audits Black Taxpayers at Higher Rates. ProPublica Investigative Report. https://www.propublica.org

Schrank, D., Eisele, B., & Lomax, T. (2021). 2021 Urban Mobility Report. Texas A&M Transportation Institute. https://mobility.tamu.edu

Smith, E., & Hendrix, A. (2022). Surtrac: Scalable Urban Traffic Control. Carnegie Mellon University Robotics Institute Technical Report.

UK HMRC. (2023). Using Data and Technology to Close the Tax Gap. HMRC Annual Report 2022–23. HM Revenue & Customs.

World Economic Forum. (2023). The Future of Government: Artificial Intelligence and the Delivery of Public Services. WEF Insight Report.