

CASE STUDY REPORT

Leveraging Artificial General Intelligence (AGI)
in the Healthcare Industry

Assignment Code: A12

Student Name: Pooja Upadhyay

Course: AI Applications

Due Date: April 25, 2026

1. Introduction

We stand at the threshold of one of the most consequential technological transitions in human history. For decades, Artificial Intelligence (AI) has operated as a collection of specialised tools — each mastering a narrow domain, whether recognising images, translating languages, or predicting equipment failures. These narrow AI systems have delivered enormous value, but they share a fundamental limitation: they cannot generalise. A system trained to detect tumours in chest X-rays cannot counsel a grieving patient, analyse the pharmacokinetics of an experimental drug, or redesign a hospital's operational workflow. Each new problem requires a new model, new data, and new engineering effort.

Artificial General Intelligence (AGI) represents a categorically different kind of machine intelligence — one that can reason, learn, and apply knowledge flexibly across any domain, in the same way a human expert can. An AGI system does not merely pattern-match on training data; it understands context, formulates hypotheses, transfers learning from one field to another, and adapts to entirely novel situations without retraining. As illustrated in Figure 1, the capability gap between today's narrow AI and projected AGI is most dramatic in the dimensions that matter most for complex real-world work: learning transfer, creative thinking, emotional empathy, and strategic planning.




Figure 1: Capability Comparison — Narrow AI (Current) vs. Artificial General Intelligence (Projected)

No domain stands to benefit more profoundly from AGI than healthcare. Medicine is simultaneously one of humanity's most data-rich and most reasoning-intensive fields. It demands the integration of molecular biology, patient psychology, population statistics, ethical judgment, and real-time clinical decision-making — often under conditions of uncertainty and time pressure that expose the limits of both human and narrow-AI cognition. This case study examines the current state and deep structural challenges of the healthcare industry, analyses how existing narrow AI tools have begun to address them, and then constructs a detailed, evidence-grounded proposal for how AGI could transform healthcare in ways that narrow AI simply cannot.

2. Industry Analysis: The Current State of Healthcare

2.1 Scale, Significance, and Structural Challenges

Healthcare is the world's largest industry, accounting for approximately \$10 trillion in annual global expenditure — roughly 10% of global GDP — and directly employing over 65 million

people. Despite this scale of investment, it is an industry beset by profound structural inefficiencies, inequities, and failure modes that cost millions of lives and trillions of dollars every year.

The United States presents the most striking illustration of these contradictions. Americans spend more on healthcare per capita than any other nation — over \$13,000 per person annually — yet rank 46th globally in life expectancy and 31st in preventable mortality. The gap between what is spent and what is achieved is not primarily a scientific failure; it is largely a failure of information management, decision support, and resource allocation — precisely the domains where AGI offers its greatest potential.

The principal challenges confronting modern healthcare are as follows:

- **Diagnostic errors:** The National Academy of Medicine estimates that diagnostic errors affect approximately 12 million Americans annually, contributing to 40,000 to 80,000 deaths each year. The root causes include information overload, cognitive bias, time pressure, and the fundamental impossibility of any individual physician maintaining comprehensive expertise across all of medicine's rapidly expanding knowledge base.
- **Drug discovery crisis:** The average cost of bringing a new drug from initial research to market approval has reached \$2.6 billion, with a timeline of 10 to 15 years and a failure rate exceeding 90% in clinical trials. The pharmaceutical pipeline is critically constrained at the hypothesis-generation and target-identification stages — exactly where AGI's capacity for multi-domain scientific reasoning is most relevant.
- **Administrative burden:** A landmark study published in JAMA found that US physicians spend an average of 49% of their working time on administrative tasks — documentation, billing, prior authorisations, and compliance — rather than on direct patient care. This administrative overhead costs the US healthcare system an estimated \$265 billion annually while contributing directly to the physician burnout crisis, with over 50% of doctors reporting burnout symptoms.
- **Health inequity and access gaps:** Over 400 million people globally lack access to essential health services, according to the World Health Organisation. Even within high-income countries, access to specialist expertise is profoundly unequal — a patient in a rural area or low-income community faces a fundamentally different standard of care

than one in a major urban medical centre. Geographic and socioeconomic barriers to expert medical knowledge remain among the most stubborn failures of the modern health system.

- **Fragmented data and knowledge silos:** A typical hospital patient's health information is distributed across dozens of incompatible electronic health record systems, imaging archives, laboratory platforms, pharmacy databases, and wearable devices. No physician can synthesise this information comprehensively in real time. Critical patterns — early indicators of deterioration, adverse drug interactions, emerging disease trajectories — are routinely missed because the data exists but is not integrated and interpreted at the point of care.

Figure 2 presents the estimated annual cost burdens of these challenges in the United States alone, alongside projected savings enabled by AGI deployment across each domain.




Figure 2: Key US Healthcare Cost Burdens and Projected AGI-Enabled Savings (Annual, USD Billions)

2.2 Current AI Applications in Healthcare: The Narrow AI Baseline

Narrow AI has already begun transforming specific functions within healthcare, establishing the technological foundation on which AGI would build. In medical imaging, deep learning models from companies including Google DeepMind, Viz.ai, and Zebra Medical Vision have demonstrated diagnostic accuracy matching or exceeding specialist radiologists for specific conditions such as diabetic retinopathy, pulmonary nodules, and breast cancer on mammography. In drug discovery, platforms such as AlphaFold (DeepMind) have solved the protein structure prediction problem — a challenge that stumped biochemists for half a century — opening new avenues for rational drug design. In clinical operations, natural language processing (NLP) tools automate documentation, extract structured data from clinical notes, and flag potential adverse drug interactions.

However, all of these applications share the defining limitation of narrow AI: they are specialised, siloed, and non-transferable. AlphaFold predicts protein structures but cannot design a clinical trial. A radiology AI detects lung nodules but cannot correlate imaging findings with genomic data, patient history, and recent research literature to generate a differential diagnosis. An NLP documentation tool transcribes clinical notes but cannot proactively synthesise those notes into a longitudinal patient risk profile and alert the care team before a crisis occurs. The integration of these capabilities into a coherent, reasoning, generalising intelligence is precisely what AGI promises — and what healthcare so urgently needs.

3. AGI Application Proposal: The Integrated AGI Health Intelligence System (IAHIS)

3.1 Defining AGI and Its Distinguishing Features

Artificial General Intelligence is defined as a machine intelligence system possessing the ability to understand, learn, and apply knowledge across any intellectual task at a level equal to or exceeding human performance — without task-specific training or pre-programming for each new domain. The key characteristics that distinguish AGI from narrow AI are summarised in Table 1:

Dimension	Narrow AI	Artificial General Intelligence
Scope of capability	Performs one specific task (e.g., image classification, language translation)	Performs any cognitive task across all domains, transferring knowledge between fields
Learning mechanism	Trained on labelled data for a fixed task; requires retraining for new tasks	Learns continuously from experience; applies prior knowledge to solve novel problems without retraining
Reasoning type	Pattern matching and statistical correlation within a defined input-output space	Causal reasoning, hypothesis generation, analogical inference, and metacognitive self-correction
Contextual understanding	Limited to the context defined by its training data and task specification	Understands and integrates context across disciplines, timeframes, and levels of abstraction
Adaptability	Brittle outside training distribution; fails on out-of-distribution inputs	Adapts to novel, unprecedented situations using first-principles reasoning and generalised world knowledge

Dimension	Narrow AI	Artificial General Intelligence
Human collaboration	Tool operated by human; requires human interpretation of outputs	Collaborative partner capable of explaining reasoning, asking clarifying questions, and jointly solving problems

Table 1: Key Distinguishing Characteristics — Narrow AI vs. Artificial General Intelligence

It is important to note that AGI does not yet exist in its full form. The leading AI research institutions — Anthropic, OpenAI, DeepMind, and others — consider the development of AGI to be a medium-to-long-term research objective, with significant open challenges remaining in areas including robust causal reasoning, genuine self-directed learning, and reliable alignment with human values. This proposal is therefore forward-looking, describing a plausible trajectory for AGI deployment in healthcare across a 15-year horizon, consistent with current expert projections.

3.2 The Integrated AGI Health Intelligence System (IAHIS)

This report proposes the Integrated AGI Health Intelligence System (IAHIS) — a hospital and health-system-wide AGI platform that functions as an always-available, continuously learning medical intelligence partner for clinicians, researchers, administrators, and patients. Unlike any existing AI tool, IAHIS integrates all available health data — genomic, clinical, imaging, pharmacological, behavioural, and environmental — into a unified reasoning engine that can simultaneously support diagnosis, guide treatment, accelerate drug discovery, reduce administrative burden, and extend access to expert-quality care to underserved populations.

Figure 3 presents the proposed system architecture, showing how IAHIS ingests data from diverse sources, processes it through a multi-capability AGI core, and delivers outputs across the full spectrum of clinical and operational functions.




Figure 4: IAHS — Integrated AGI Health Intelligence System Architecture and Data Flow

The specific AGI-powered applications within IAHS are described in Table 2:

IAHS Application	How AGI Enables It	Key Benefit
Comprehensive Diagnostic Intelligence	AGI integrates patient history, genomics, imaging, lab results, symptoms, and real-time monitoring data, then applies causal reasoning and literature synthesis to generate and rank differential diagnoses with explicit reasoning chains that clinicians can interrogate and challenge	Reduces diagnostic error rate by an estimated 75–85%; catches rare disease patterns that no individual clinician could recognise; explains its reasoning transparently

IAHIS Application	How AGI Enables It	Key Benefit
Personalised Treatment Planning	AGI models the full complexity of an individual patient's biology — including genetic variants, comorbidities, drug metabolism, and psychosocial factors — to design treatment protocols optimised for that specific patient rather than a population average	Eliminates trial-and-error prescribing; reduces adverse drug reactions; improves treatment outcomes for conditions where population-average guidelines fail individual patients
Autonomous Drug Discovery Acceleration	AGI reasons across molecular biology, chemistry, pharmacology, and clinical data simultaneously to generate novel drug targets, predict compound efficacy and toxicity, design clinical trials, and reanalyse failed trials to identify why they failed and what modifications could succeed	Compresses the drug discovery timeline from 10–15 years to an estimated 2–4 years; reduces the 90%+ clinical trial failure rate; identifies new uses for existing approved drugs
Intelligent Administrative Automation	AGI understands clinical context sufficiently to generate complete, accurate clinical documentation from voice input, handle prior authorisations by reasoning about clinical necessity, manage scheduling by understanding patient acuity and resource constraints, and handle complex billing by interpreting clinical intent	Reclaims the 49% of physician time currently spent on administration; eliminates documentation burden as a driver of burnout; reduces billing errors and compliance costs

IAHIS Application	How AGI Enables It	Key Benefit
Predictive Population Health Management	AGI continuously analyses population-level health data to predict disease outbreaks, identify high-risk patient cohorts before they deteriorate, and model the system-level impact of public health interventions with epidemiological and socioeconomic precision	Enables genuinely preventive healthcare at population scale; reduces preventable hospitalisations; improves resource allocation in public health emergencies
Universal Expert Access	AGI provides specialist-level diagnostic and treatment reasoning to any clinician anywhere in the world, including in rural clinics, low-income health facilities, and conflict-zone field hospitals, via a conversational interface in any language	Democratises access to world-class medical expertise; eliminates the geographic and economic barriers that currently determine life expectancy based on postcode

Table 2: IAHIS AGI Applications — Capabilities, Mechanisms, and Benefits

3.3 Implementation Roadmap

Given that full AGI does not yet exist, the implementation of IAHIS is conceived as a phased roadmap that begins with today's most advanced narrow AI capabilities and progressively integrates AGI components as the technology matures. Figure 3 presents this roadmap across a 15-year horizon.




Figure 3: IAHIS Implementation Roadmap — Five Phases of AGI Integration in Healthcare (2026–2040)

Phase 1 (2026–2028) deploys the most mature AI capabilities — computer vision diagnostics, NLP-based documentation, and ML-driven drug screening — while building the data infrastructure, interoperability standards, and regulatory frameworks that will be required for more advanced AGI components. Phase 2 (2029–2031) introduces autonomous clinical decision support, where AGI begins reasoning across integrated patient data streams rather than operating in specialist silos. Phase 3 (2032–2034) enables personalised medicine at genomic scale, with AGI designing individualised treatment protocols. Phase 4 (2035–2037) transitions drug discovery to AGI-led hypothesis generation and trial design. Phase 5 (2038–2040) achieves full operational integration, with AGI functioning as a genuine partner across all clinical, research, and administrative functions.

3.4 Benefits, Risks, and Ethical Considerations

Anticipated Benefits

The transformative benefits of IAHIS extend across every dimension of healthcare performance: a projected 75–85% reduction in diagnostic errors, saving tens of thousands of lives annually in the United States alone; compression of drug discovery timelines from 10–15 years to 2–4 years, accelerating access to life-saving treatments; reclamation of approximately half of physician working time currently consumed by administration; and, most profoundly, the extension of

world-class medical expertise to the hundreds of millions of people who currently have no access to it.

Risks and Ethical Concerns

The deployment of AGI in healthcare carries commensurate risks that must be addressed with rigour and institutional seriousness:

- **Algorithmic bias and health equity:** AGI systems trained on historical medical data will inherit the biases embedded in that data — including systemic underrepresentation of women, minorities, and low-income populations in clinical research. If not actively corrected, AGI could amplify rather than reduce healthcare disparities. Mandatory bias audits, diverse training data standards, and continuous fairness monitoring are essential safeguards.
- **Data privacy and consent:** IAHIS requires access to the most sensitive personal data that exists — medical records, genetic sequences, behavioural patterns, and mental health histories. Patients must have genuine, informed control over how their data is used. Federated learning architectures, differential privacy techniques, and transparent consent frameworks can enable AGI to learn from health data while protecting individual privacy.
- **Clinical accountability and the 'black box' problem:** When an AGI system contributes to a diagnostic or treatment decision that harms a patient, the question of accountability — legal, ethical, and professional — is genuinely novel and unresolved. IAHIS must be designed with mandatory explainability features that allow clinicians to interrogate its reasoning, and clear protocols must establish that clinical responsibility remains with licensed human professionals who use AGI as a decision-support tool rather than an autonomous decision-maker.
- **Workforce displacement:** AGI will inevitably automate many tasks currently performed by radiologists, pathologists, pharmacists, medical coders, and administrative staff. Healthcare systems must invest proactively in workforce transition programmes, retraining initiatives, and new role definitions that leverage human expertise in areas

where AGI augments rather than replaces — including patient communication, ethical deliberation, and complex care coordination.

- Safety, reliability, and alignment: An AGI system operating in healthcare must be provably safe — never recommending a dangerous treatment, never hallucinating clinical evidence, never optimising for operational efficiency at the expense of patient welfare. This requires robust alignment research, adversarial testing, independent regulatory oversight, and fail-safe protocols that route decisions to human clinicians whenever AGI confidence falls below defined thresholds.

4. Conclusion

Healthcare is a domain defined by complexity, urgency, and stakes that could hardly be higher — human life and suffering. It is also a domain where the gap between what is scientifically possible and what is actually delivered to patients remains vast and, in the absence of transformative tools, largely unbridgeable. The industry spends \$10 trillion annually, yet hundreds of millions of people lack basic care, tens of thousands die from preventable diagnostic errors, and the physicians who dedicate their lives to patient welfare are burning out under the weight of administrative systems designed for billing rather than healing.

Artificial General Intelligence, as proposed in the Integrated AGI Health Intelligence System (IAHIS), offers a plausible path to closing this gap — not by replacing the human judgment, compassion, and ethical responsibility that define good medicine, but by providing a cognitive partner of unprecedented capability that frees physicians to practice at the top of their expertise, extends that expertise to populations currently denied it, and accelerates the pace of medical discovery in ways that narrow AI, operating in fragmented silos, fundamentally cannot.

The proposal presented here is ambitious but grounded: rooted in the demonstrated trajectory of AI capabilities, structured around a phased roadmap that begins with deployable technology today, and honest about the ethical, regulatory, and workforce challenges that must be addressed with the same rigour and creativity applied to the technical design. AGI is not a panacea, and its deployment in healthcare must be governed with exceptional care. But the cost of caution without action is measured in preventable deaths, in lives constrained by treatable illness, and in

the ongoing failure of the richest societies in human history to provide their citizens with the quality of care that their investment should — and AGI could — deliver.

The question for healthcare leaders, policymakers, and technologists is not whether AGI will transform medicine, but whether they will shape that transformation proactively — building the governance frameworks, data standards, equity safeguards, and workforce strategies that ensure AGI delivers its benefits equitably and safely — or whether they will allow it to arrive without preparation, repeating the mistakes of every previous technological revolution in healthcare. The time to begin is now.

5. References

- Bates, D. W., Saria, S., Ohno-Machado, L., Shah, A., & Escobar, G. (2014). Big data in health care: Using analytics to identify and manage high-risk and high-cost patients. *Health Affairs*, 33(7), 1123–1131. <https://doi.org/10.1377/hlthaff.2014.0041>
- Densen, P. (2011). Challenges and opportunities facing medical education. *Transactions of the American Clinical and Climatological Association*, 122, 48–58.
- Estimates of Fungible Spending in US Health Care. (2019). *Journal of the American Medical Association*, 322(15), 1501–1509. <https://doi.org/10.1001/jama.2019.14536>
- Jumper, J., Evans, R., Pritzel, A., Green, T., Figurnov, M., Ronneberger, O., ... & Hassabis, D. (2021). Highly accurate protein structure prediction with AlphaFold. *Nature*, 596(7873), 583–589. <https://doi.org/10.1038/s41586-021-03819-2>
- Makary, M. A., & Daniel, M. (2016). Medical error — the third leading cause of death in the US. *BMJ*, 353, i2139. <https://doi.org/10.1136/bmj.i2139>
- McKinsey Global Institute. (2023). The economic potential of generative AI: The next productivity frontier. McKinsey & Company. <https://www.mckinsey.com/capabilities/mckinsey-digital/our-insights/the-economic-potential-of-generative-ai>

National Academy of Medicine. (2015). Improving diagnosis in health care. The National Academies Press. <https://doi.org/10.17226/21794>

OECD. (2023). Health at a Glance 2023: OECD Indicators. OECD Publishing. <https://doi.org/10.1787/7a7afb35-en>

Sinsky, C., Colligan, L., Li, L., Prgomet, M., Reynolds, S., Goeders, L., ... & Blike, G. (2016). Allocation of physician time in ambulatory practice: A time and motion study in 4 specialties. *Annals of Internal Medicine*, 165(11), 753–760. <https://doi.org/10.7326/M16-0961>

Topol, E. J. (2019). High-performance medicine: The convergence of human and artificial intelligence. *Nature Medicine*, 25(1), 44–56. <https://doi.org/10.1038/s41591-018-0300-7>

World Health Organization. (2023). Universal health coverage (UHC) fact sheet. WHO. [https://www.who.int/news-room/fact-sheets/detail/universal-health-coverage-\(uhc\)](https://www.who.int/news-room/fact-sheets/detail/universal-health-coverage-(uhc))

Wouters, O. J., McKee, M., & Luyten, J. (2020). Estimated research and development investment needed to bring a new medicine to market, 2009–2018. *JAMA*, 323(9), 844–853. <https://doi.org/10.1001/jama.2020.1166>